

# M2PGi – Internet Of Things

Pr. Olivier Gruber

(olivier.gruber@univ-grenoble-alpes.fr)

Laboratoire d'Informatique de Grenoble  
Université de Grenoble-Alpes

# Goal... Bare-metal Discovery...

- Widening your horizon as a developer
  - What you probably know about...



Probably application programming, right?

C, Java, JavaScript, Python...

Basic usage of your operating system, right?

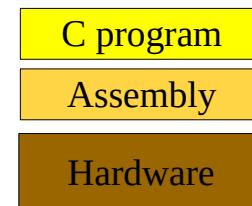
File system and Graphical User Interface  
Maybe a hint of shell usage/scripting...

- What you probably know less about...



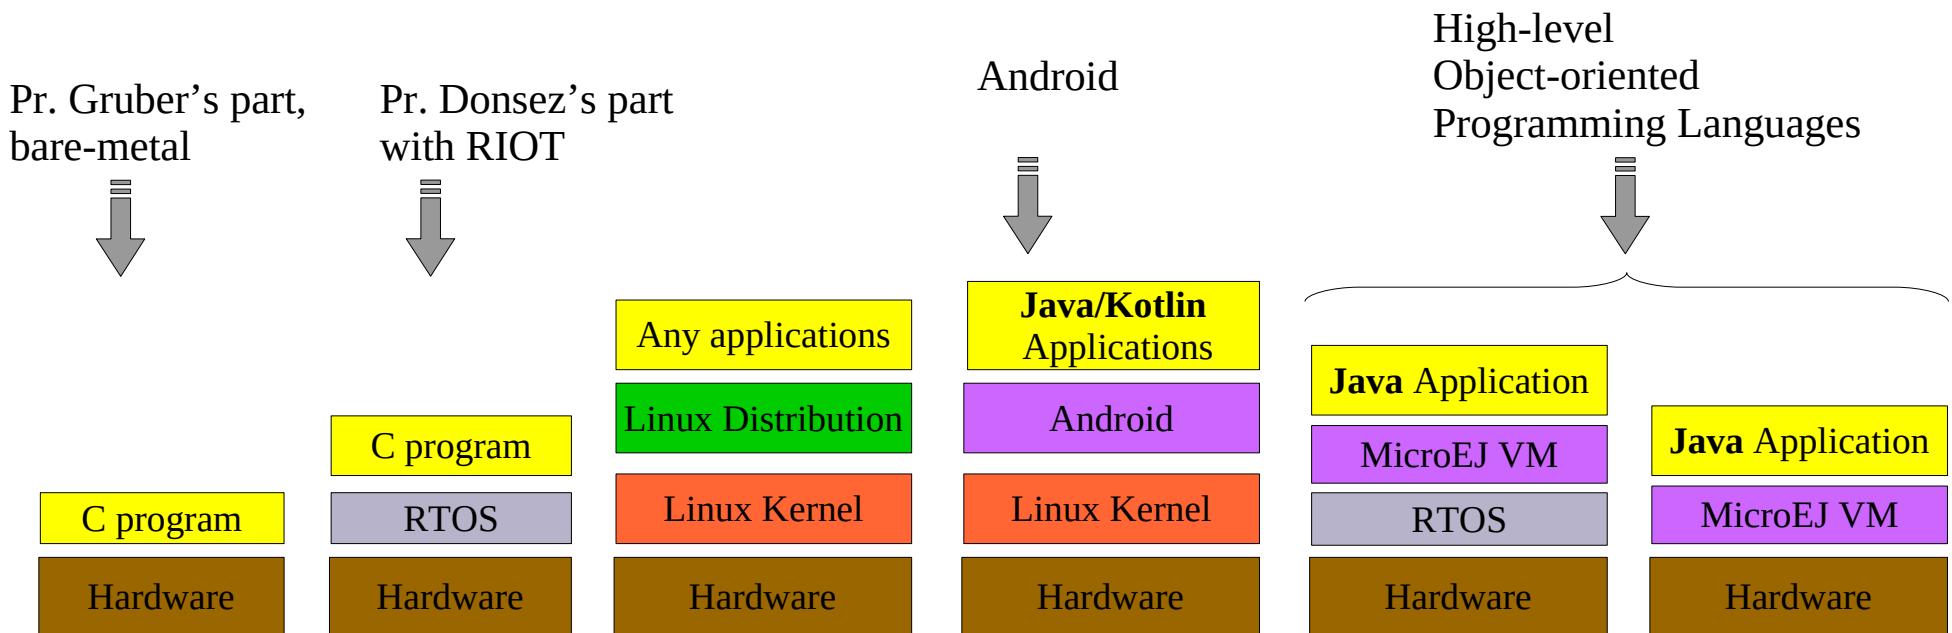
**How do we program these?**

**Which tools do we use?**



**Bare-metal  
programming**

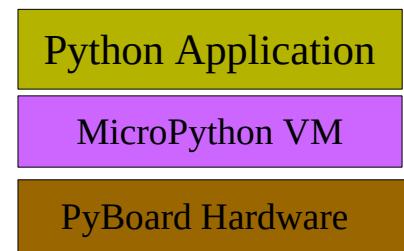
# Embedded Development Spectrum



- STM32F405RG microcontroller
  - **168 MHz Cortex M4 CPU with hardware floating point**
  - **1024 KB flash ROM**
  - **192 KB RAM**
  - Micro USB connector for power and serial communication
  - Micro SD card slot
  - GPIOs

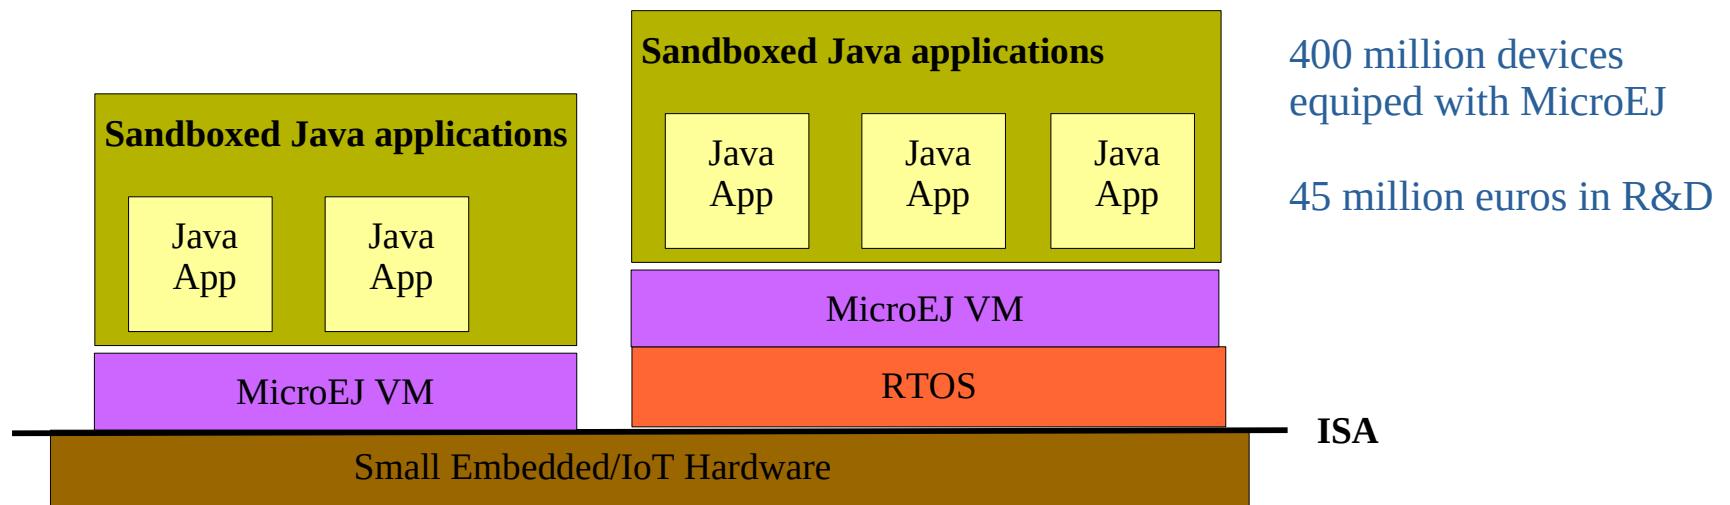


Micro-Python





- Micro Embedded Java
  - Java for small and smart embedded devices
  - Similar virtualization technology as Android but *much tighter* implementation
    - Android: \$15 processor, 32MB of memory
    - MicroEJ: \$1 processor, 128KB of memory
  - Google Cloud IoT Partners
    - IoT solutions that leverage Google's secure, global, and scalable infrastructure



(1) <https://www.microej.com/>

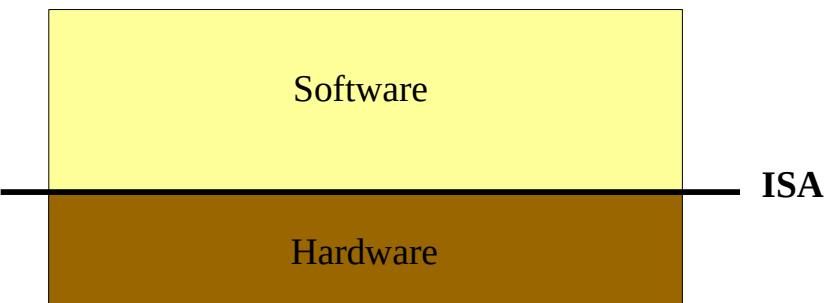
# This Part Pedagogy

---

- Team learning, individual work
  - Hands-on learning and coding bare-metal software
  - Your own **work-log: tracking what you do and what you learned**
  - This work-log document is first and foremost for your own records
- The destination is not the goal, the goal is the learning along that path
  - Be curious... expand your skills and know-how
  - Discuss what you learned/understood
  - Explain to others – Ask questions
  - **WRITE DOWN WHAT YOU HAVE LEARNED / WHAT YOU DID**
- Part-I Evaluation
  - No exam – but weekly progress checkpoints
  - **Every week:** you will surrender your work (**work-log** and your **code**)
  - Your weekly involvement is the **largest part of your final grade**

# Bare-metal Development

- Bare-metal Software Basics
  - Runs directly on the "bare metal"
- Instruction Set Architecture (ISA)
  - Defines the instruction set (user and privileged)
  - Defines other privileged concepts (interrupts, traps, page tables, ...)
- Cross-development Toolchain
  - Compiler, linker, debugger, and other tools



# Bare Metal – Development Environment

- Using a real board
  - Relies on using a USB-Serial cable
  - For both JTAG and a console (/dev/ttyUSB0)
- JTAG through /dev/ttyUSB0
  - JTAG to upload the firmware
  - JTAG for hardware and software debugging
- Using a terminal emulator on your laptop
  - Terminal emulator like minicom or kermit
  - **Serial line for a console** (a command-line interface)
    - Send characters for printing on the screen
    - Receive characters from the keyboard



USB – Serial Cable – RS 232  
JTAG debugging  
Console over serial Line

A screenshot of a terminal window titled "Terminal emulator on /dev/ttyUSB0". The window shows the kernel boot process, including CPU initialization, memory mapping, and support for various chipsets. A red arrow points from the text "Terminal" to this window.

```
o.gruber@pingouin: ~ 79x32
[ 0.000000] Initializing cgroup subsys cpuset
[ 0.000000] Initializing cgroup subsys cpu
[ 0.000000] Initializing cgroup subsys cpacct
[ 0.000000] Linux version 3.13.0-106-generic (buildd@lcy01-30) (gcc version 4.8.4 (Ubuntu 4.8.4-2ubuntu1-14.04.3) ) #153-Ubuntu SMP Tue Dec 6 15:44:32 UTC 2016 (Ubuntu 3.13.0-106.153-generic 3.13.11-ckt39)
[ 0.000000] Command line: BOOT_IMAGE=/boot/vmlinuz-3.13.0-106-generic root=UUID=cdf1fcf-57a4-4b96-a476-a90b40b72bbb ro quiet splash vt.handoff=7
[ 0.000000] KERNEL supported cpus:
[ 0.000000]   Intel GenuineIntel
[ 0.000000]   AMD AuthenticAMD
[ 0.000000]   Centaur CentaurHauls
[ 0.000000] e820: BIOS-provided physical RAM map:
[ 0.000000]   BIOS-e820: [mem 0x0000000000000000-0x0000000000000cff] usable
```

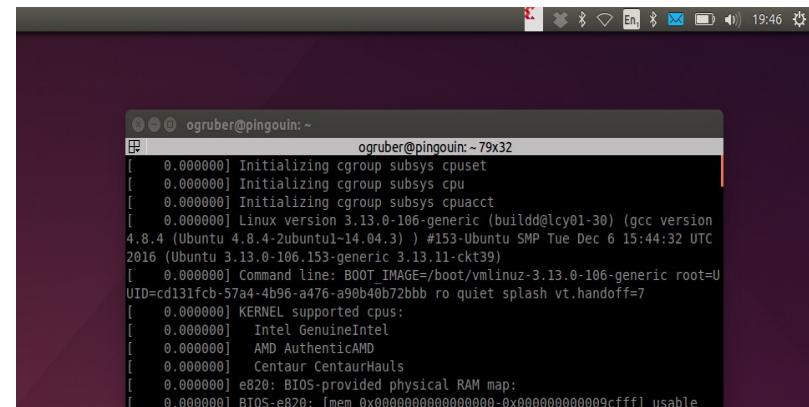


Terminal

# Bare Metal – Development Environment

- Using an emulator – QEMU

- A regular Linux process
- Emulates a machine for your bare-metal software
- Direct support for GDB debugging
- Serial line for a command-line interface



```
ogruber@pingouin:~ 79x32
[ 0.00000] Initializing cgroup subsys cpuset
[ 0.00000] Initializing cgroup subsys cpu
[ 0.00000] Initializing cgroup subsys cputacct
[ 0.00000] Linux version 3.13.0-106-generic (buildd@lcy01-30) (gcc version 4.8.4 (Ubuntu 4.8.4-2ubuntu1-14.04.3) ) #153-Ubuntu SMP Tue Dec 6 15:44:32 UTC 2016 (Ubuntu 3.13.0-106.153-generic 3.13.11-ckt39)
[ 0.00000] Command line: BOOT_IMAGE=/boot/vmlinuz-3.13.0-106-generic root=UUID=cd131fcf-57a4-4b96-a476-a90b40b72bbb ro quiet splash vt.handoff=7
[ 0.00000] KERNEL supported cpus:
[ 0.00000]   Intel GenuineIntel
[ 0.00000]   AMD AuthenticAMD
[ 0.00000]   Centaur CentaurHauls
[ 0.00000] e820: BIOS-provided physical RAM map:
[ 0.00000]   BIOS-e820: [mem 0x0000000000000000-0x000000000009cff] usable
```

"The matrix is a prison that you cannot see, taste, or smell."  
*Morpheus*



Board



Serial Line



Terminal



Terminal



*For you: an all-software-development experience in the comfort of your chair!*



USB – Serial Cable – RS 232

# Emulation with QEMU – Board and Terminal

9

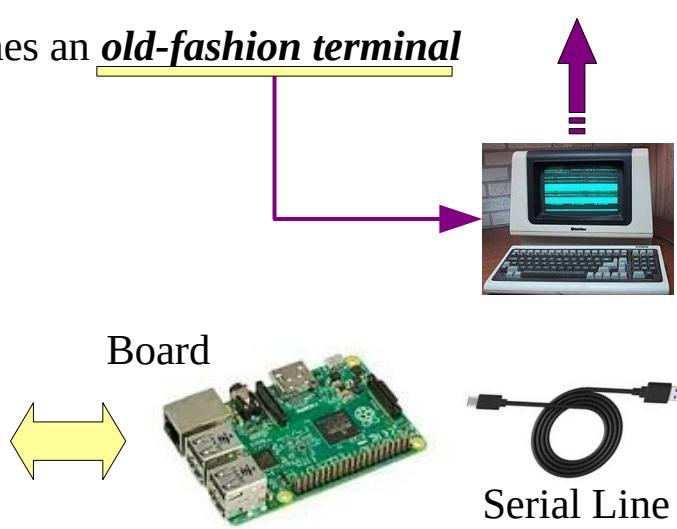
## In one terminal:

- **cross-compile** and **link** your kernel
- **launch QEMU**, it will load your kernel

```
$ arm-none-eabi-gcc -g -o kernel.elf *.o  
$ qemu-system-arm ... -device loader,file=kernel.elf
```

That terminal becomes an **old-fashion terminal**

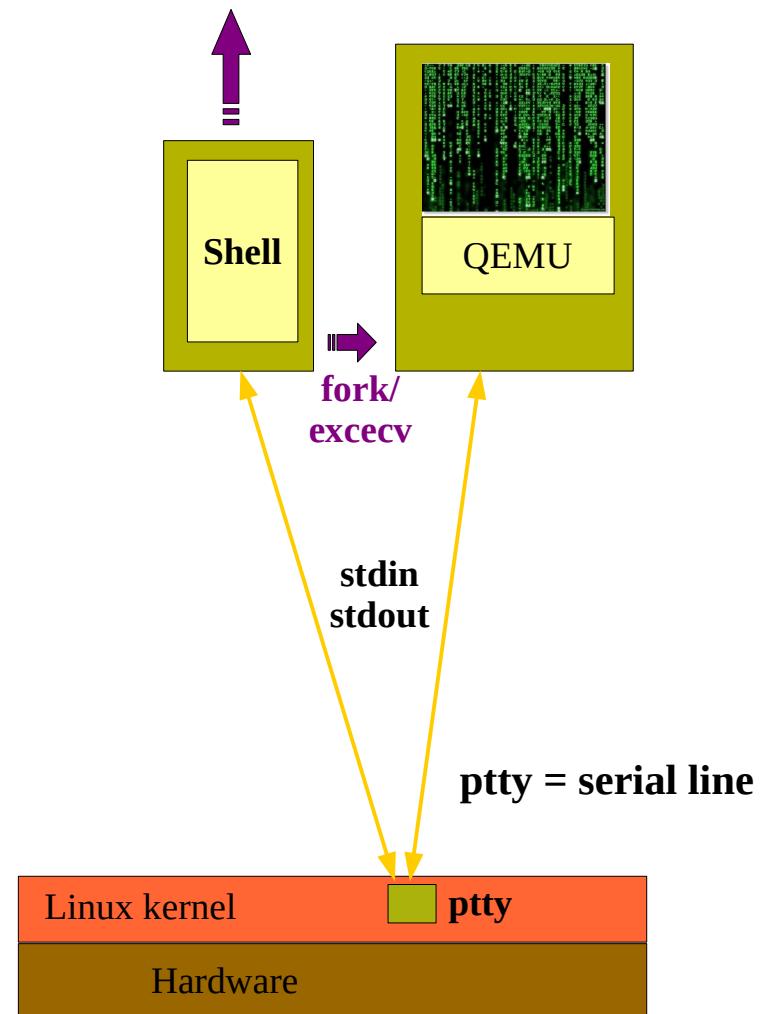
- **no more echo**
- **no more edit**
- **cursor anywhere**



Sending characters  
through the serial line  
will print them on the terminal



Received characters  
will be the one typed  
on the keyboard



# Asking QEMU to Stop the Illusion!



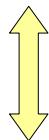
10

Open the QEMU  
console and quit  
the emulation

**Ctrl-A c**

```
$ qemu-system-arm -M versatilepb -cpu cortex-a8 -m "32K"  
-nographic -serial mon:stdio -device loader,file=kernel.elf  
QEMU 8.2.2 monitor - type 'help' for more information  
...  
(qemu) quit  
$
```

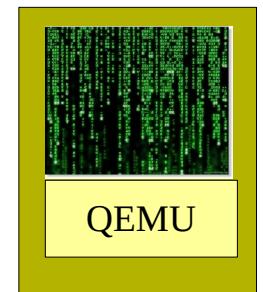
Or simply use **Ctrl-A x**



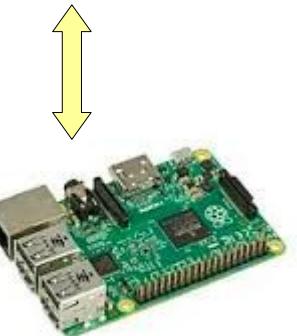
Terminal



Serial Line



QEMU



VersatilePB

# Debugging with QEMU

11

## In one terminal:

- cross-compile and link your kernel
- launch QEMU to load your kernel

```
$ arm-none-eabi-gcc -g -o kernel.elf *.o  
$ qemu-system-arm ... -gdb tcp::1234 -S
```



## In another terminal:

- launch your ARM debugger
- with the ELF file of your kernel
- but attach to a remote target

```
$ arm-none-eabi-gdb kernel.elf  
...  
(gdb) target remote localhost:1234
```

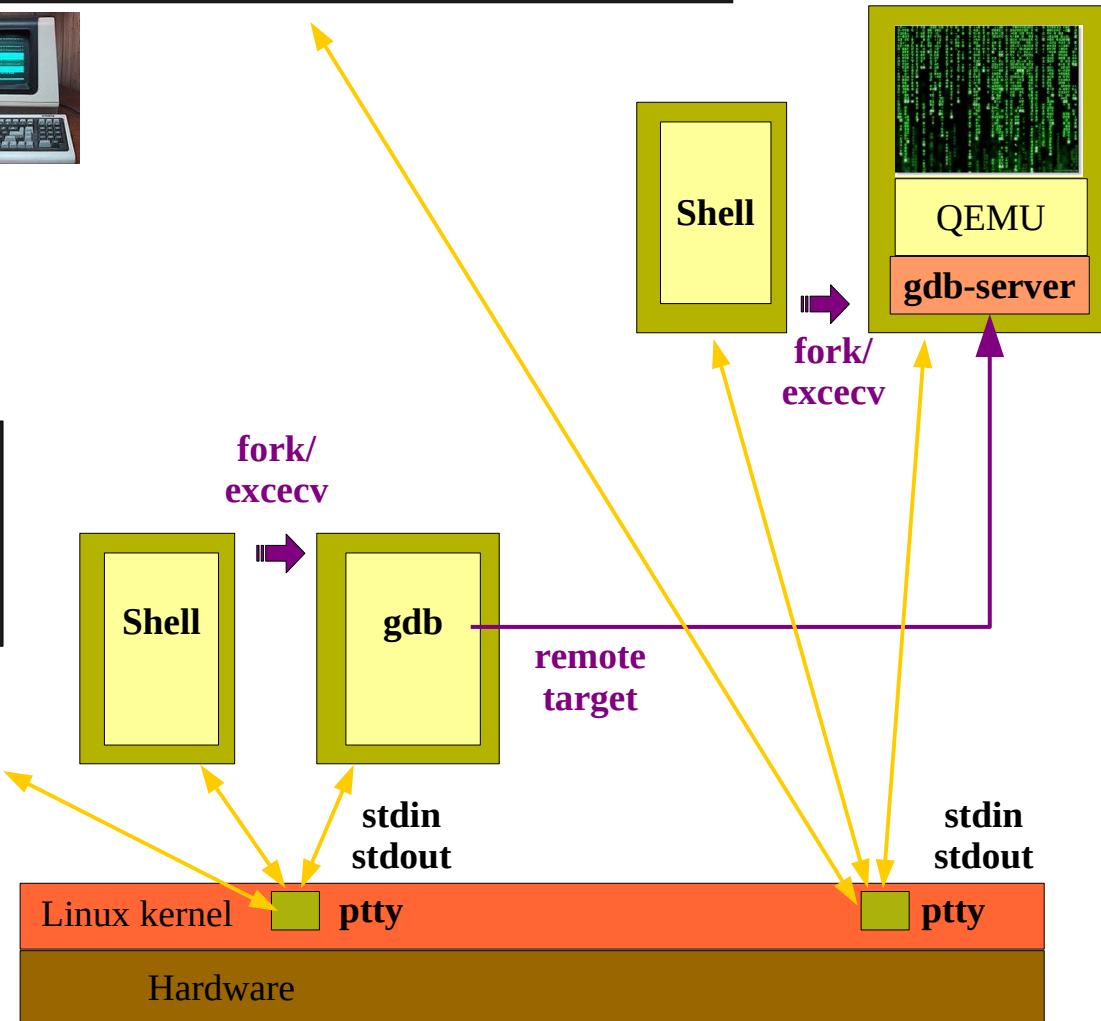
### Nota Bene:

you should use the short notation

**tar rem :1234**

instead of

**target remote localhost:1234**



# Debugging with QEMU

12

```
$ qemu-system-arm -M versatilepb -cpu cortex-a8 -m "16K"  
-nographic -serial mon:stdio -device loader,file=kernel.elf -gdb tcp::1234 -S
```

```
$ gdb-multiarch kernel.elf  
gdb-multiarch kernel.elf  
GNU gdb (Ubuntu 15.0.50.20240403-0ubuntu1) 15.0.50.20240403-git  
Copyright (C) 2024 Free Software Foundation, Inc.  
  
(gdb) tar rem :1234  
Remote debugging using :1234  
?? () at exception.s:31  
31      ldr pc, reset_handler_addr  
(gdb) print /x $r15  
$1 = 0x0  
(gdb)
```

Where is the execution stopped at?

- At address 0x0000-0000
- Register **r15** is the “program counter”

Why is it stopped there?

What code is there to execute?

How did we build that code?

# Linker Script – Code / Data Memory Layout

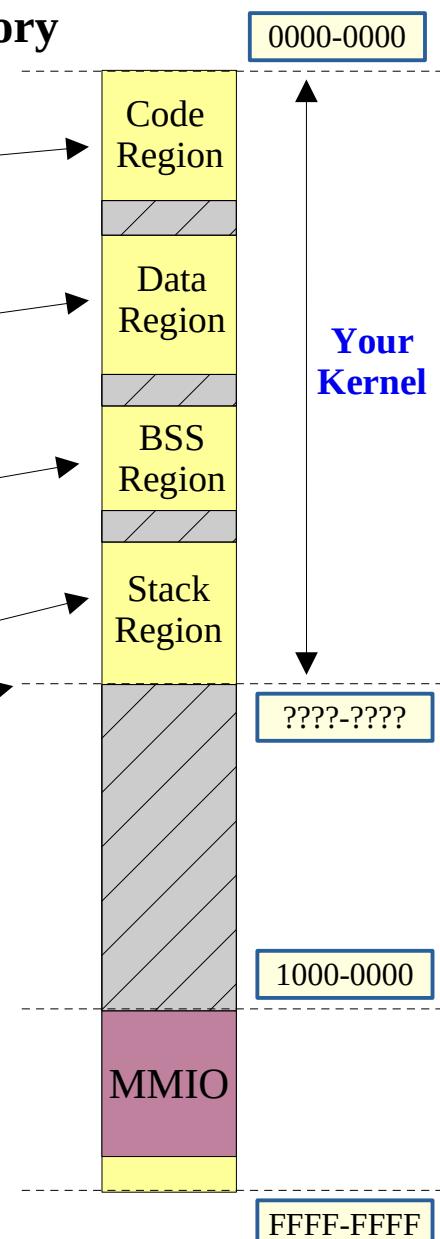
13

## Linker Script given to GCC: “kernel.ld”

```
SECTIONS
{
    . = 0x00;
    .text : {
        build/exception.o(.text)
        build/startup.o(.text)
        build/*(.text)
    }
    . = ALIGN(4);
    .data : {
        build/*(.data)
    }
    . = ALIGN(4);
    _bss_start = .;
    .bss . : {
        build/*(.bss COMMON)
    }
    . = ALIGN(4);
    _bss_end = .;

    . = ALIGN(8);
    . = . + 0x1000;
    stack_top = .;
}
```

## Board Memory



QEMU: respect sections when loading “elf files”

```
$ qemu-system-arm -M versatilepb -cpu cortex-a8 -m "32K"
-nographic -serial mon:stdio -device loader,file=kernel.elf
```

# Cross-compilation & QEMU-based Execution

```
MACHINE=versatilepb
CPU=cortex-a8
MEMSIZE_IN_KB=16

QEMU=qemu-system-arm
QEMU_ARGS= -M $(MACHINE) -cpu $(CPU) -M $(MEMSIZE_IN_KB)K -nographic -serial mon:stdio
CFLAGS= -g -DMEMORY="($(MEMSIZE_IN_KB)*1024)"
ASFLAGS= -g
CFLAGS+= -mcpu=$(CPU) -c -nostdlib -ffreestanding
LDFLAGS= -g -mcpu=$(CPU) -T kernel.ld -nostdlib -static

# Object files to build and link together
#-----
OBJS= startup.o main.o exception.o

# Compilation Rules
#-----
%.o: %.c
    arm-none-eabi-gcc $(CFLAGS) -o $@ $<
%.o: %.s
    arm-none-eabi-as $(ASFLAGS) -o $@ $<

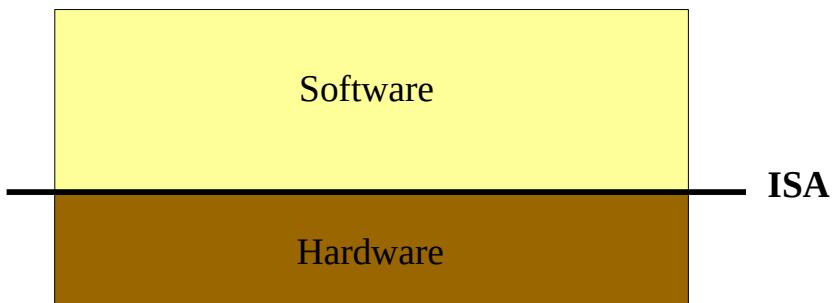
# Build and link all, producing an ELF and binary
#-----
all: $(OBJS)
    arm-none-eabi-ld $(LDFLAGS) $^ -o kernel.elf
    arm-none-eabi-objcopy -O binary kernel.elf kernel.bin

run: all
    $(QEMU) $(QEMU_ARGS) -device loader,file=kernel.elf

debug: all
    $(QEMU) $(QEMU_ARGS) -device loader,file=kernel.elf -gdb tcp::1234 -S
```

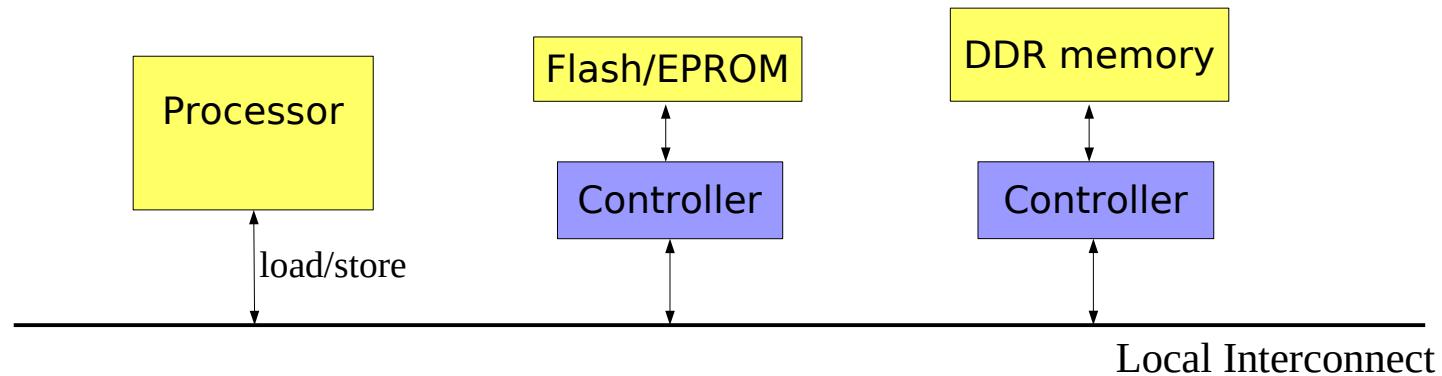
# Bare-metal Development Basics

- Bare-metal Software
  - Runs directly on the "bare metal"
- Instruction Set Architecture (ISA)
  - Defines the instruction set (user and privileged)
  - Defines other privileged concepts (page tables, traps, interrupts, ...)
- Cross-development Toolchain
  - Compiler, linker, debugger, and other tools
- Using Integrated Development Environments (IDEs)

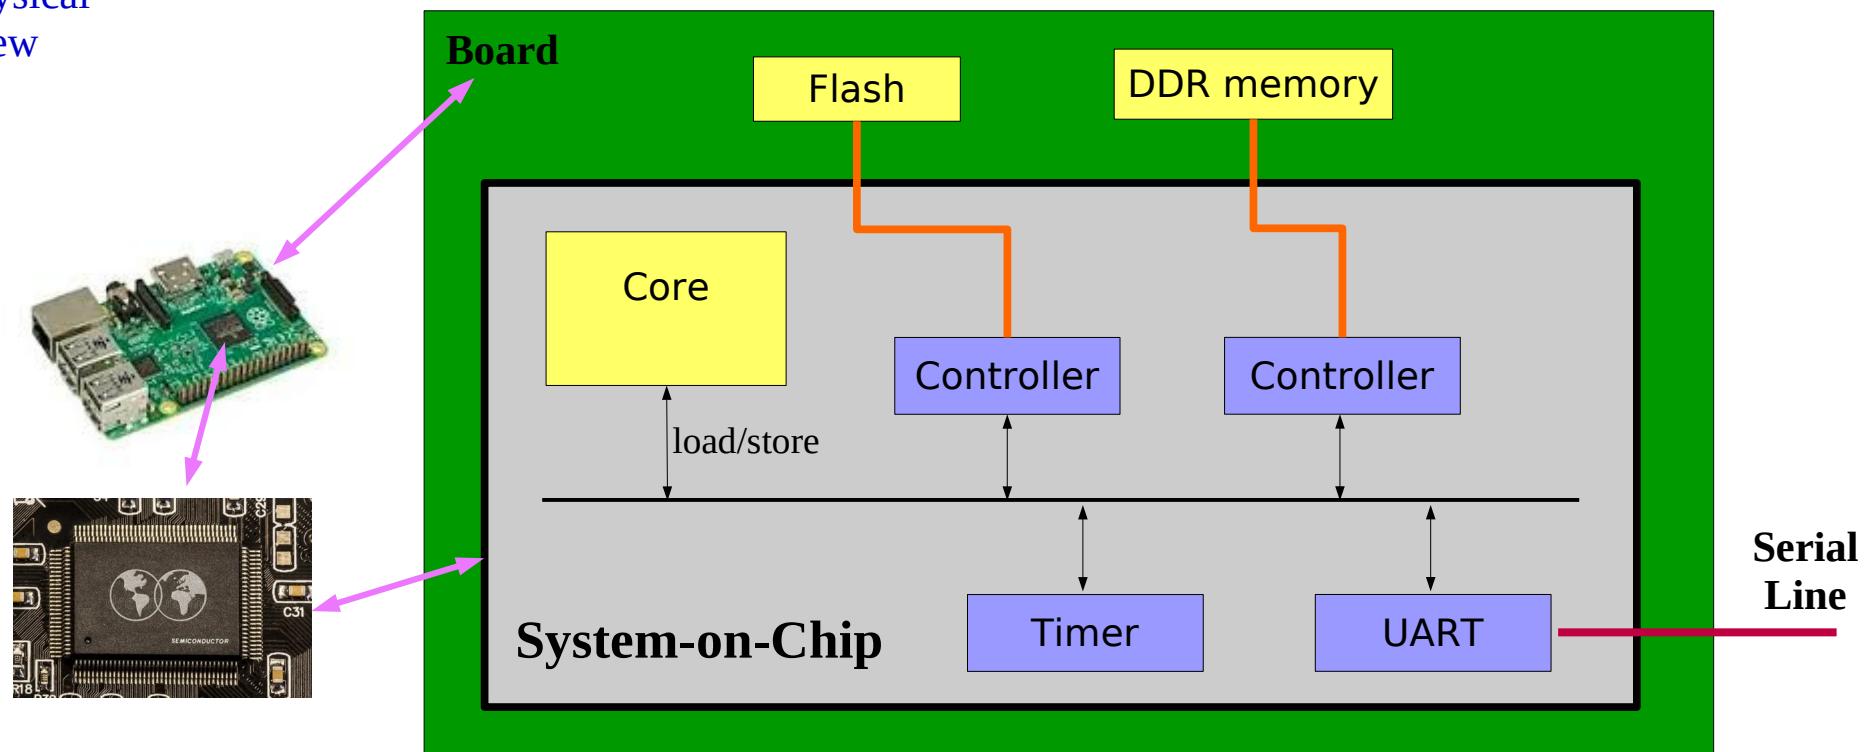


# Hardware – Basics

Conceptual View

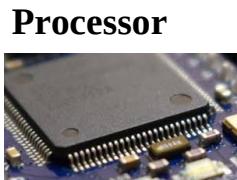
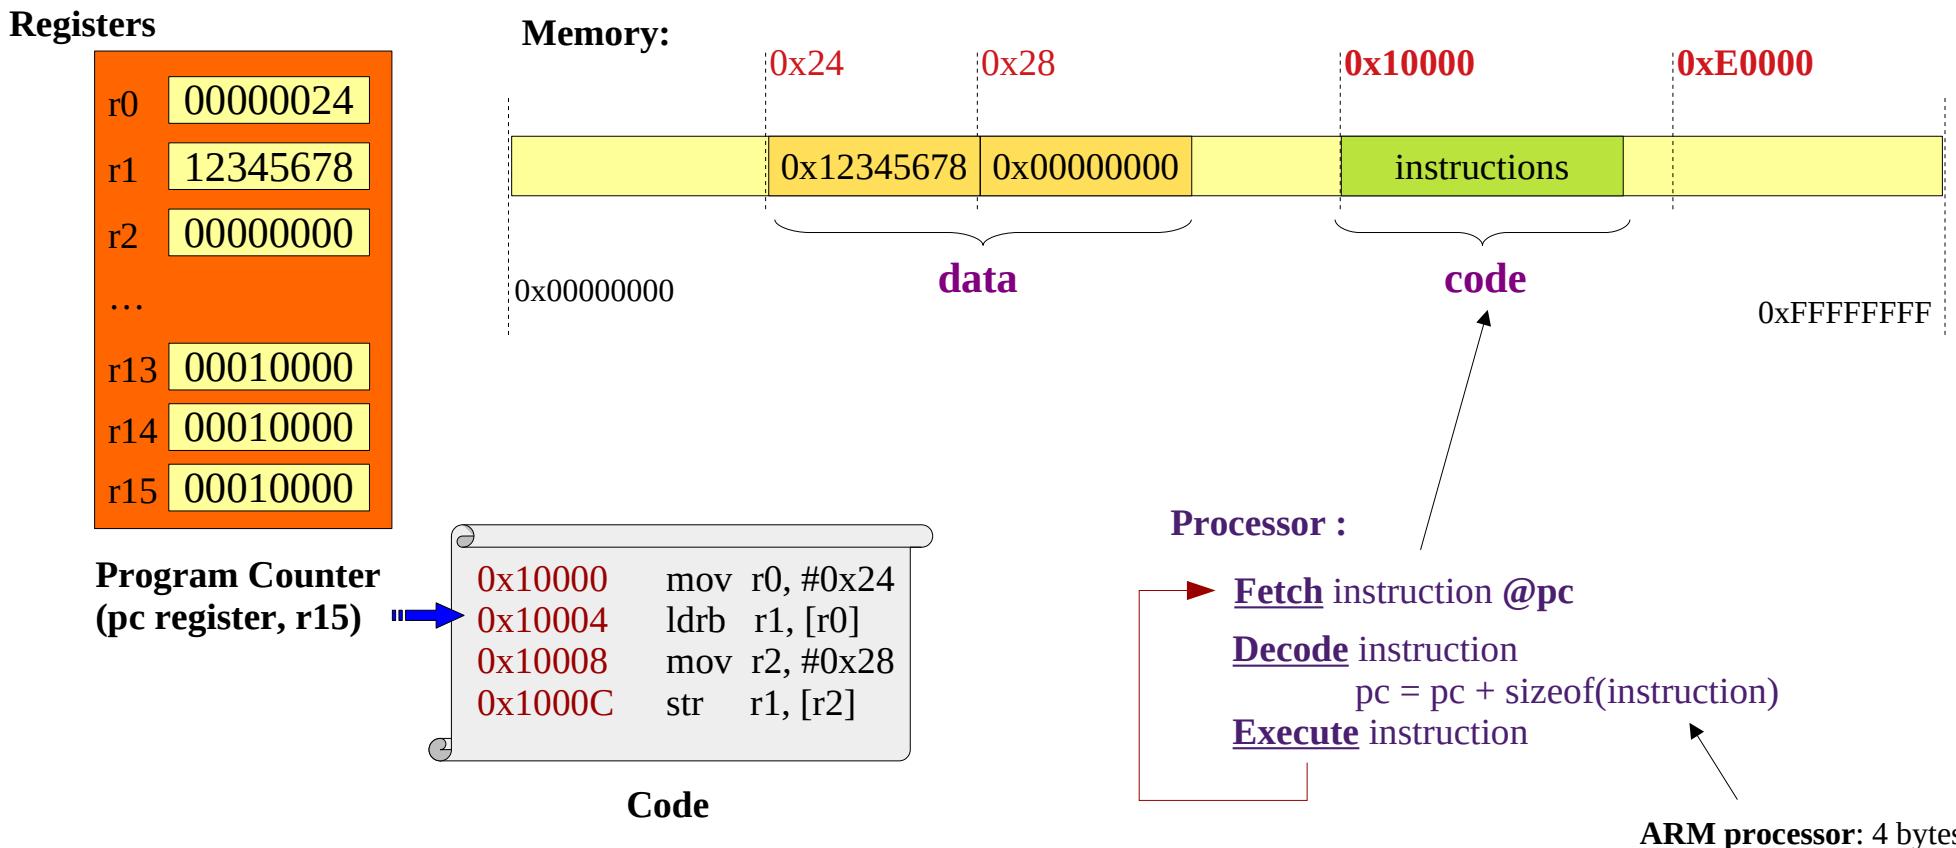


Physical View



# Reminder – Processor & Memory

17



**Processor**

Only two operations:  
- value ← **load**(address)  
- **store**(address,value)  
both for data and code...

**Memory**



**Controller**

**Bus/Interconnect**

# Example: Zynq-7000 – Boot Sequence

18

**Processor Reset**  
→ pc=0x0 000-0000

**With:**  
**OCM is mapped low**

**Boot Steps:**

1. initialize DDR
2. copy boot code  
above 0x0000-FFFF
3. jump there
4. ...

Table 4-1: System-Level Address Map

Address Range	CPUs and ACP	AXI_HP	Other Bus Masters <sup>(1)</sup>	Notes
0000-0000 to 0000-FFFF	OCM	OCM	OCM	Address not filtered by SCU and OCM is mapped low
0000_0000 to 0003_FFFF <sup>(2)</sup>	DDR	OCM	OCM	Address filtered by SCU and OCM is mapped low
	DDR			Address filtered by SCU and OCM is not mapped low
				Address not filtered by SCU and OCM is not mapped low
				Address filtered by SCU
0004_0000 to 0007_FFFF	DDR			Address not filtered by SCU
0008_0000 to 000F_FFFF	DDR	DDR	DDR	Address filtered by SCU
0010_0000 to 3FFF_FFFF	DDR	DDR	DDR	Address not filtered by SCU <sup>(3)</sup>
4000_0000 to 7FFF_FFFF	PL		PL	General Purpose Port #0 to the PL, M_AXI_GP0
8000_0000 to BFFF_FFFF	PL		PL	General Purpose Port #1 to the PL, M_AXI_GP1
E000_0000 to E02F_FFFF	IOP		IOP	I/O Peripheral registers, see <a href="#">Table 4-6</a>
E100_0000 to E5FF_FFFF	SMC		SMC	SMC Memories, see <a href="#">Table 4-5</a>
F800_0000 to F800_0BFF	SLCR		SLCR	SLCR registers, see <a href="#">Table 4-3</a>
F800_1000 to F880_FFFF	PS		PS	PS System registers, see <a href="#">Table 4-7</a>
F890_0000 to F8F0_2FFF	CPU			CPU Private registers, see <a href="#">Table 4-4</a>
FC00_0000 to FDFF_FFFF <sup>(4)</sup>	Quad-SPI		Quad-SPI	Quad-SPI linear address for linear mode
FFFC_0000 to FFFF_FFFF <sup>(2)</sup>	OCM	OCM	OCM	OCM is mapped high
				OCM is not mapped high



# Example: Zynq-7000 – Boot Sequence

19

**Memory Map:**  
Normal memory (DDR)  
from 0x0000-00000  
to 0x3FFF-FFFF

**Boot Steps:**

1. initialize DDR
2. copy boot code  
above 0x0000-FFFF
3. jump there
4. **map OCM high**
5. initialize boot device,  
like SD-card reader
6. **Load more code/data...**

Table 4-1: System-Level Address Map

Address Range	CPUs and ACP	AXI_HP	Other Bus Masters <sup>(1)</sup>	Notes
	OCM	OCM	OCM	Address not filtered by SCU and OCM is mapped low
0000_0000 to 0003_FFFF <sup>(2)</sup>	DDR	OCM	OCM	Address filtered by SCU and OCM is mapped low
	DDR			Address filtered by SCU and OCM is not mapped low
0004_0000 to 0007_FFFF	DDR			Address not filtered by SCU and OCM is not mapped low
0008_0000 to 000F_FFFF	DDR	DDR	DDR	Address filtered by SCU
0010_0000 to 3FFF_FFFF	DDR	DDR	DDR	Address not filtered by SCU <sup>(3)</sup>
4000_0000 to 7FFF_FFFF	PL		PL	General Purpose Port #0 to the PL, M_AXI_GP0
8000_0000 to BFFF_FFFF	PL		PL	General Purpose Port #1 to the PL, M_AXI_GP1
E000_0000 to E02F_FFFF	IOP		IOP	I/O Peripheral registers, see <a href="#">Table 4-6</a>
E100_0000 to E5FF_FFFF	SMC		SMC	SMC Memories, see <a href="#">Table 4-5</a>
F800_0000 to F800_0BFF	SLCR		SLCR	SLCR registers, see <a href="#">Table 4-3</a>
F800_1000 to F880_FFFF	PS		PS	PS System registers, see <a href="#">Table 4-7</a>
F890_0000 to F8F0_2FFF	CPU			CPU Private registers, see <a href="#">Table 4-4</a>
FC00_0000 to FDFF_FFFF <sup>(4)</sup>	Quad-SPI		Quad-SPI	Quad-SPI linear address for linear mode
FFFC_0000 to FFFF_FFFF <sup>(2)</sup>	OCM	OCM	OCM	OCM is mapped high
				OCM is not mapped high

# WARNING

20

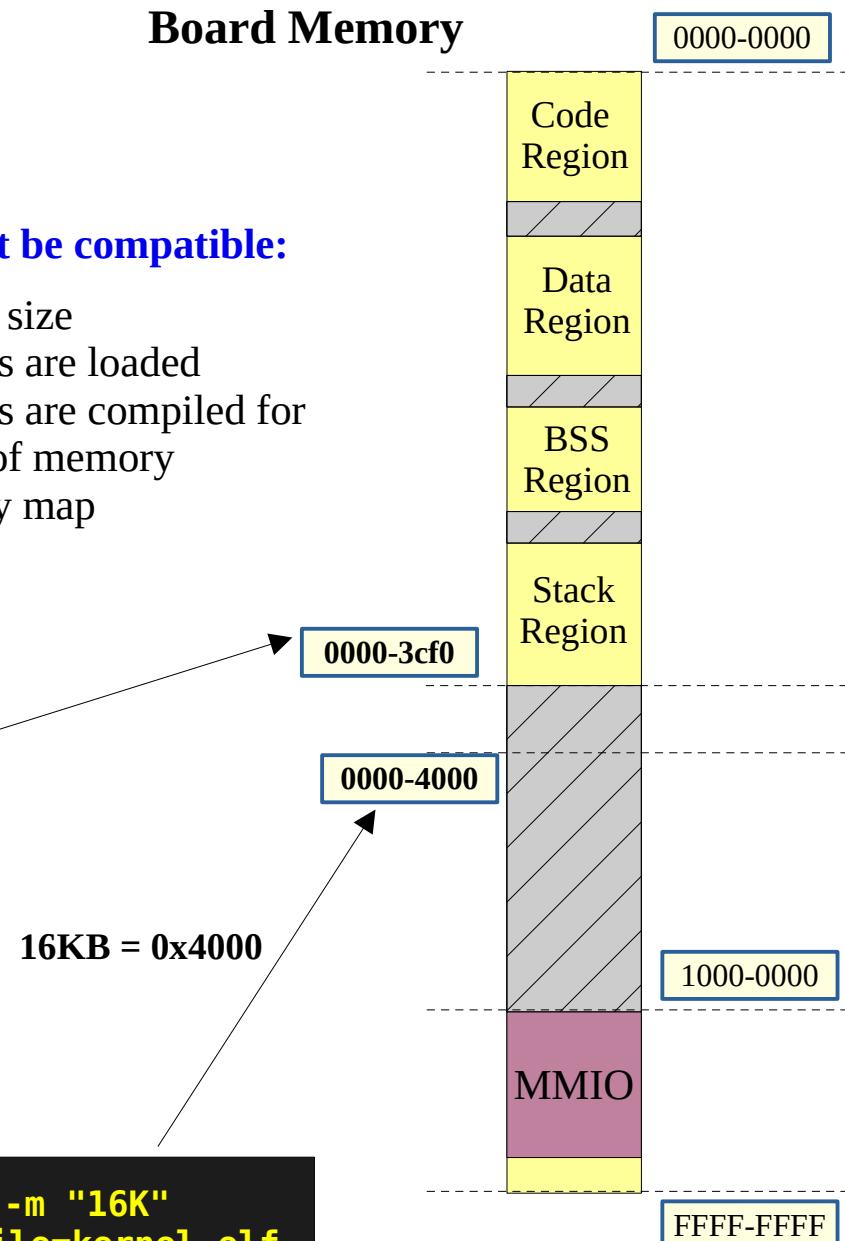
```
SECTIONS
{
    . = 0x00;
    .text : {
        build/exception.o(.text)
        build/startup.o(.text)
        build/*(.text)
    }
    . = ALIGN(4);
    .data : {
        build/*(.data)
    }
    . = ALIGN(4);
    _bss_start = .;
    .bss . : {
        build/*(.bss COMMON)
    }
    . = ALIGN(4);
    _bss_end = .;

    . = ALIGN(8);
    . = . + 0x1000;
stack_top = .;
}
```

## Everything must be compatible:

- code and data size
- where sections are loaded
- where sections are compiled for
- total amount of memory
- board memory map

**stack\_top**



QEMU: respect sections when loading “elf files”

```
$ qemu-system-arm -M versatilepb -cpu cortex-a8 -m "16K"
  -nographic -serial mon:stdio -device loader,file=kernel.elf
```

# Cortex-A8 Exception Vector & Code (“exception.s”)

```
// Hardware Exception Vector (loaded 0x0000-0000, see linker script)

0x00    ldr pc, reset_handler_addr
0x04    ldr pc, undef_handler_addr
0x08    ldr pc, swi_handler_addr
0x0C    ldr pc, prefetch_abort_handler_addr
0x10    ldr pc, data_abort_handler_addr
0x14    ldr pc, unused_addr
0x18    ldr pc, irq_handler_addr
0x1C    ldr pc, fiq_handler_addr

0x20  reset_handler_addr: .word _reset_handler
0x24  undef_handler_addr: .word _undef_handler
0x28  swi_handler_addr: .word _swi_handler
0x2C  prefetch_abort_handler_addr: .word _prefetch_abort_handler
0x30  data_abort_handler_addr: .word _data_abort_handler
0x34  not_used_addr: .word _not_used_handler
0x38  irq_handler_addr: .word _isr_handler
0x3C  fiq_handler_addr: .word _fiq_handler

0x40  _isr_handler:           // unexpected interrupt
0x44    b _isr_handler
0x48  _undef_handler:        // undefined instruction
0x4C    b _unused_handler
0x50  _fiq_handler:          // unexpected fast interrupt
0x54    b _fiq_handler
0x58  _undef_handler:        // unexpected undefined instruction
0x5C    b _undef_handler
0x60  _swi_handler:          // unexpected software interrupt
0x64    b _swi_handler
0x68  _prefetch_abort_handler: // unexpected prefetch-abort
0x6C    b _prefetch_abort_handler
0x70  _data_abort_handler:   // unexpected abort
0x74    b _data_abort_handler

_reset_handler:
...   // unexpected abort trap
```

QEMU starts the execution at 0x0000-0000...

It is the “reset exception”...

ldr pc, reset\_handler\_addr  
→ ldr pc, [pc+0x18]

because pipeline of depth two,  
so two instructions,  
so 0x08 bytes.

# Cortex-A8 – Assembly Boot Code (“startup.s”)

```
.global _reset_handler
_reset_handler:

/*
 * Set the core in the SYS_MODE, with all interrupts disabled.
 */
msr    cpsr_c, #(CPSR_SYS_MODE | CPSR_IRQ_FLAG | CPSR_FIQ_FLAG)

/* set the C stack pointer for the SYS_MODE */
ldr    sp, =stack_top

/*
 * Clear out the bss section, located from _bss_start to _bss_end.
 * This is a C convention, the GNU GCC compiler will group
 * all global variables that need to be zeroed on startup
 * in the bss section of the ELF.
 */
ldr    r4, =_bss_start
ldr    r9, =_bss_end
mov    r5, #0

1:
stmia  r4!, {r5}
cmp    r4, r9
blo    1b

/*
 * Set the GCC frame-pointer to null (marks the begin of the stack)
 * and then upcall the C entry function _start(void)
 */
eor    r11, r11, r11
ldr    r3, =_start
blx   r3

halt:
    b halt // in case the function "_start" returns...
```

Setup the processor mode,  
and disable all interrupts  
Then set the stack pointer.

Then clear the "bss data"

Done with assembly,  
for now, upcalling C code...

# Cortex-A8 Exception Vector & Code (“exception.s”)

```
// Hardware Exception Vector (loaded 0x0000-0000, see linker script)

ldr pc, reset_handler_addr
ldr pc, undef_handler_addr
ldr pc, swi_handler_addr
ldr pc, prefetch_abort_handler_addr
ldr pc, data_abort_handler_addr
ldr pc, unused_addr
ldr pc, irq_handler_addr
ldr pc, fiq_handler_addr

reset_handler_addr: .word _reset_handler
undef_handler_addr: .word _undef_handler
swi_handler_addr: .word _swi_handler
prefetch_abort_handler_addr: .word _prefetch_abort_handler
data_abort_handler_addr: .word _data_abort_handler
not_used_addr: .word _not_used_handler
irq_handler_addr: .word _isr_handler
fiq_handler_addr: .word _fiq_handler

_undef_handler: // undefined instruction
...
1: b 1b // infinite loop for debugging

_prefetch_abort_handler: // prefetch-abort trap
1: b 1b // infinite loop for debugging

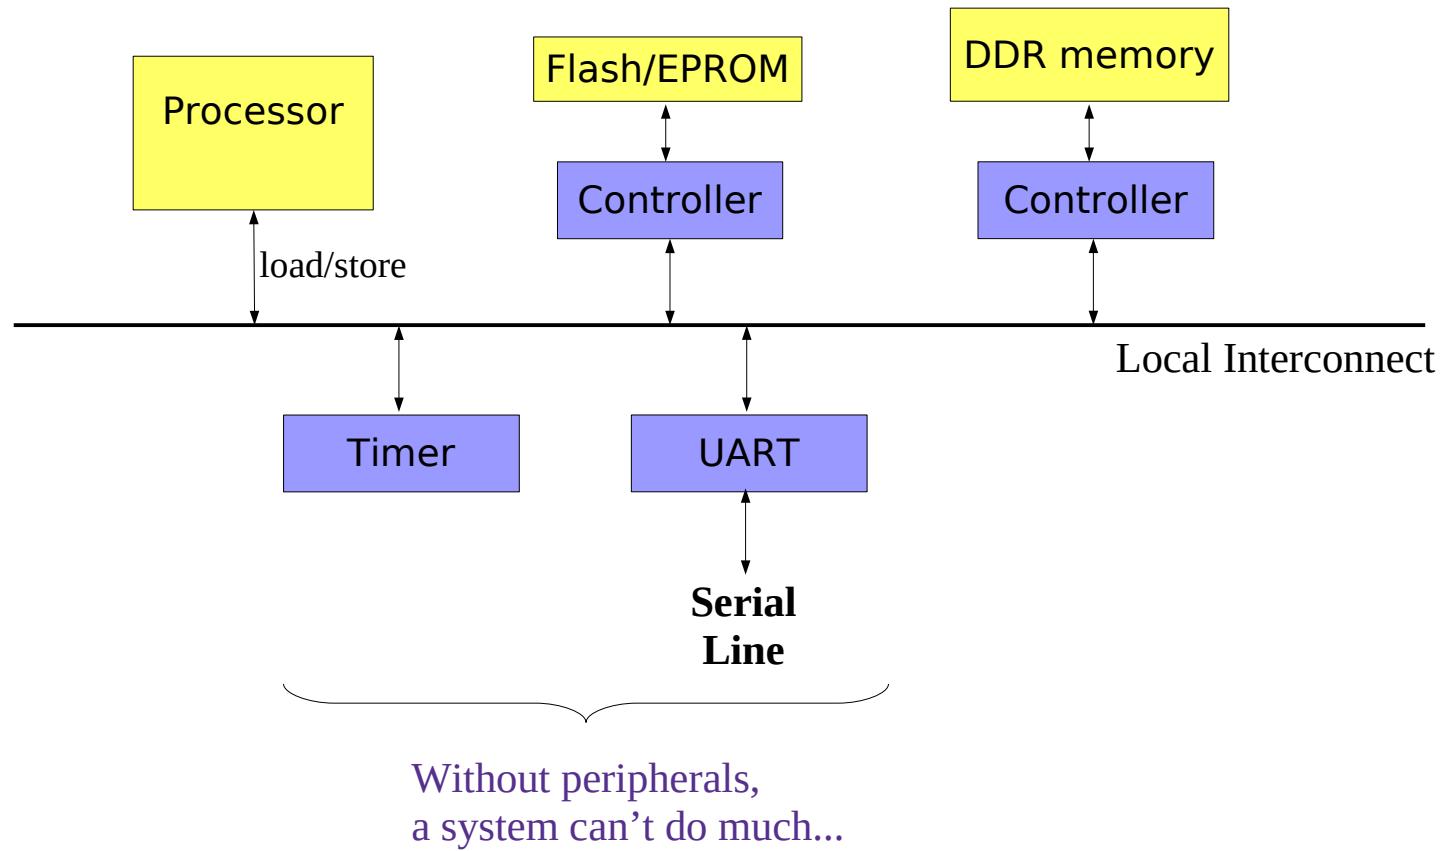
_data_abort_handler: // load/store abort trap
...
1: b 1b // infinite loop for debugging
```

Prefetch-abort and data-abort hardware exceptions are triggered by the memory subsystem.

The undef hardware exception is triggered by the core itself, when it tries to execute an undefined instruction

The execution will trap here when your program has bugs...

# Next Time – Hardware Peripherals



Terminal



Serial Line



VersatilePB